

PROPOSAL

A governed legal-skills pipeline for Harrow & Vale LLP

Deterministically-constrained Claude skills — starting with term-sheet review and due-diligence — that a lawyer will actually reach for mid-deal.

Prepared for: Tom Harrow, Ops & Knowledge Lead (for sign-off by Priya Vale, Managing Partner)

Engagement: Hackathon 1 — Synthetic Signal Associate Programme

Date: 22 July 2026 **Check-in / demo:** Wednesday 29 July 2026

Budget envelope: £2,000–£5,000 (Phase 1, fixed fee)

Prepared following discovery calls with Tom Harrow and review of the firm's data-room samples (four term-sheet formats, Priya's fixed DD checklist, and the GreenGrid cap table & Articles).

Contents

1. Executive summary

Every lawyer at Harrow & Vale already uses Claude. The problem isn't adoption — it's that ten people are each reinventing their own prompts and skills in isolation, nothing is shared or kept current, and nobody has a clean answer to the question the partners ask before anything gets signed off: where does our client data actually go?

We propose one thing, not ten: **a single governed pipeline where any skill a lawyer builds is contract-bound and deterministically eval-gated before it is approved, versioned, and installed firm-wide.** We prove it by putting the firm's two bread-and-butter tasks through the gate first — Priya's term-sheet review and the "Thursday afternoon dump" due-diligence check — and by handing you a one-page, partner-ready answer on data handling.

This is deliberately phased. Phase 1 is a fixed-fee proof-of-concept sized to sit inside your £2,000–£5,000 envelope: two working skills, demonstrated running on your own samples on 29 July, plus the data memo. The firm-wide pipeline and governance (Phase 2) is only worth building once Phase 1 has proved it saves real hours — which is exactly the order Tom asked for.

A working prototype already exists. Across all four of your term-sheet formats, an eval-driven second iteration of the term-sheet skill moved from an overall reliability score of 0.50 to 1.00, with a regression gate that blocks any future change from shipping if it makes the skill worse.

2. What we heard

From discovery with Tom, and from the data-room samples, three things are non-negotiable:

- **Priya wants an Exception Report, not a book report.** She reads term sheets herself. She wants economics first (valuation and basis, liquidation preference, option pool), then control (board, investor vetoes, drag-along) — mapped against the firm's BVCA-aligned baseline, flagging where a term is aggressive and what is missing entirely, as actionable bullets.
- **The checklist is fixed.** Priya's due-diligence checklist is applied to every matter and must be used verbatim. A tool respects it; it does not reinvent its own categories.
- **Data handling is the real gate.** The blocker to sign-off is not the budget — it is the partners asking where client data goes. You are on a Claude Team plan, documents live in Microsoft 365 / SharePoint, and one associate occasionally uses GitHub. The answer has to speak to exactly that.

Tom's own trigger made the value concrete: *the "Thursday afternoon dump"* — thirty due-diligence documents arrive at 4:30pm, feedback due Friday noon, and today that means tired associates manually mapping files against the checklist. The tool has to turn that into an instant "14 satisfied, 6 missing, here's the document and page for each."

3. The solution: one governed legal-skills pipeline

The deliverable Harrow & Vale keeps is not a clever prompt — it is a gate. Anything a lawyer builds passes through the same three steps before it is trusted on a live deal:

Step	What it means	Why it makes it "kosher"
1. Contract	Every skill declares a strict JSON-schema output. Critical extraction is typed — never free-form prose.	The output shape is fixed and machine-checkable: deterministically constrained.
2. Prove	Each skill ships a golden dataset of known-good and known-bad examples. An adversarial evaluator scores precision & recall and fails the build on regression.	You can always answer "why do we trust this skill?" with numbers, not vibes.

Step	What it means	Why it makes it "kosher"
3. Ship	Passing skills publish to the firm's private repo with a version and changelog. Lawyers install from one place; v2 arrives automatically.	One shared, current, vetted home — the opposite of ten private copies.

3.1 Skill one — the term-sheet Exception Report

Point the skill at any term sheet and it returns Priya's report: economics first, then control, mapped against the BVCA-aligned baseline, with each finding tagged **aggressive**, **watch**, or **info**, plus a list of what's missing — every item carrying a citation and a confidence score. It handles your four formats (SAFE, priced round, convertible note, and the terse seed summary) by extracting on meaning, not layout. A worked example is in Appendix A.

3.2 Skill two — the due-diligence checklist mapper

Drop a folder of due-diligence documents and the skill maps them against Priya's fixed checklist verbatim, returning each item as satisfied, partial, or missing — with the document name and location for the ones it found, and a chase-list for the ones it didn't. On the GreenGrid sample set the prototype reports "2 satisfied, 1 partial, 25 missing of 28", each with a citation. That is Tom's Thursday-afternoon problem, handled before he's finished reading the covering email.

3.3 How we prove it works — the eval harness

A working skill is only the baseline. What separates a student project from something a partner will sign off is proving the skill behaves **consistently across every format** and doesn't quietly regress when someone tweaks a prompt. The prototype already does this:

Term-sheet skill	v1 (first pass)	v2 (after eval-driven iteration)	
Instrument detection	0.75	1.00	across 4 formats
Exception detection (F1)	0.29	1.00	penalises misses AND false alarms
Missing-item detection (F1)	0.50	1.00	
Overall reliability	0.496	1.000	gate threshold 0.90 → PASS

The evaluator is deliberately strict about *false alarms* as well as misses — a tool that cries wolf is one a lawyer stops trusting. The regression gate means every future change is measured before it ships.

3.4 The pipeline and governance (Phase 2)

Once the two skills have earned their place, the same gate generalises to the whole firm. The governance model is simple and administrator-friendly:

- **Author:** any lawyer or associate writes the skill, its contract, and a handful of golden examples.
- **Approver:** Tom signs the pipeline run for ops/security; Priya owns the domain baselines (the BVCA baseline, the DD checklist) that skills must respect verbatim.
- **Promotion:** a change re-runs the harness in CI. Green and no regression → version bumped, published, auto-served to every install. Red → blocked at the current version.
- **Audit:** every approved version keeps its eval report, so the firm can always show its working.

4. Where your data goes — the sign-off answer

This is the question the partners will ask first, so here it is in plain terms, specific to your stack (Claude Team plan; documents in Microsoft 365 / SharePoint; a private GitHub repo). *This is a starting memo for Priya — the specifics should be confirmed against Anthropic's current Commercial Terms and Data Processing Addendum before sign-off.*

- **Training:** Under Anthropic's Commercial Terms, inputs and outputs from commercial products (Claude Team / "Claude for Work") are not used to train Anthropic's models by default. The consumer-terms changes around extended retention and training do not apply to Team/Enterprise.
- **Retention:** Conversations are retained in the product so the team has continuity, and can be deleted from the dashboard. Zero-data-retention arrangements exist but are enabled per-organisation for qualifying Enterprise accounts via Anthropic's account team — worth raising if the partners want no retention at all.
- **Where documents live:** Client documents stay in your existing M365 / SharePoint governance. The skills operate on a document in-session; they do not create a new data store. A term sheet is only sent to Claude at the moment a lawyer runs the review.
- **The GitHub repo holds code, not client data:** the private repo stores skill definitions, contracts, and golden datasets. Golden examples must use synthetic or redacted content (as the hackathon samples do). A firm rule — never commit a client document to the repo — keeps the confidentiality boundary clean.
- **Recommendation:** for a firm handling this data, evaluate Claude Enterprise for SSO, audit logs, longer/stricter admin controls and ZDR eligibility. Team is fine to prove Phase 1; Enterprise is the likely home for firm-wide rollout.

Full memo with citations delivered as a one-pager for Priya. Primary sources: Anthropic Privacy Center and Commercial Terms (see footnotes in the memo).

5. Scope, effort and price

Priced as a fixed-fee proof-of-concept, not a day rate. We've sized Phase 1 to sit inside your stated envelope on purpose — the aim is to earn Phase 2 by saving you real hours, not to maximise the first invoice. Two options so you can choose the scope:

Phase 1 option	What's included	Effort	Fixed fee
A — Recommended	Both skills (term-sheet Exception Report + DD checklist mapper), the eval harness & regression gate, and the one-page data memo. Demoed running on your samples on 29 July.	~9–10 days	£4,500
B — Leaner	Term-sheet Exception Report + eval harness + data memo. DD mapper deferred to Phase 2.	~6 days	£2,800

Phase 2 (indicative, scoped after Phase 1): firm-wide pipeline, the versioning/approval process, install path for all ten lawyers, and additional skills as the firm identifies them. Indicative range £8,000–£15,000 depending on how many skills and whether you move to Enterprise. Not billed now — and not worth committing to until Phase 1 has proved itself.

6. What success looks like

We'll judge Phase 1 the way Tom framed it — would a lawyer actually reach for this mid-deal?

- A lawyer runs the term-sheet skill on a live sheet and gets Priya's Exception Report in the shape she wants, consistently, across formats.
- Tom runs the DD mapper on a document dump and gets an accurate satisfied/missing report with citations, in seconds rather than an evening.
- Priya signs off because the data question has a real, written answer.

- The eval harness shows the skills clear the reliability threshold with no regression — the number the firm can point to.

7. Next steps

- **You pick Option A or B** (or tell us to adjust scope).
- **A 20-minute call** to confirm the BVCA baseline reflects the firm's house view and to walk the DD checklist edge cases. (Tom offered his direct line for this.)
- **We demo both skills running on your samples at the 29 July check-in**, hand over the data memo for Priya, and agree whether to proceed to Phase 2.

Thank you — we're ready to start on your word.

Appendix A — Sample Exception Report output

Two of your four samples, as the skill returns them (abridged from the structured JSON). Note the economics-then-control order and the aggressive / watch tags.

GreenGrid Analytics — Series A (priced round)

- **[watch] Option pool** — No option pool created in this round; future hires dilute post-round. Consider a pre-money pool. (Cap-table note.) Confidence 0.82.
- **[missing] Pro-rata rights** — No pro-rata/participation right stated; unusual to omit for a lead. Confidence 0.80.
- **[info] 1x non-participating preference, broad-based weighted-average anti-dilution, 2 founder / 1 investor board** — all standard. Figures reconcile with the cap table and Articles.

Anchorline Biotech — Convertible loan note

- **[aggressive] Change-of-control premium** — 1.5x premium on a change of control before conversion; materially investor-favourable vs a 1x baseline. Confidence 0.92.
- **[watch] Interest** — 8% per annum, compounding annually — above the 0–6% simple norm; increases the conversion/redemption amount. Confidence 0.88.
- **[missing] Pro-rata rights** — None stated for the note holder. Confidence 0.76.

Appendix B — Eval harness output (extract)

```
=== Skill version v2 ===
nimbus-safe      instrument=OK | exc P/R/F1=1.00/1.00/1.00 | missing F1=1.00
greengrid-series-a instrument=OK | exc P/R/F1=1.00/1.00/1.00 | missing F1=1.00
anchorline-convertible instrument=OK | exc P/R/F1=1.00/1.00/1.00 | missing F1=1.00
solace-seed      instrument=OK | exc P/R/F1=1.00/1.00/1.00 | missing F1=1.00
AGG instrument=1.00 exc_F1=1.00 miss_F1=1.00 OVERALL=1.000
=== REGRESSION GATE === v1 0.496 -> v2 1.000 RESULT: PASS
```

Full prototype (skills, contract schema, golden datasets, evaluator, and the DD mapper) accompanies this proposal as a runnable folder.